

The Live App Walkthrough

This document mirrors the running application. Every sample question below shows the app's **actual output, quoted verbatim** — the routing line, the trust badge, the headline, the metric tiles, and the warnings — plus what to point at and what to say. Learn this, and you can narrate the demo perfectly even if you had to present from paper.

Meridian Pipeline Intelligence · outputs captured from the deterministic engine (AI narration adds phrasing on top, never new numbers) · Sai Saran

Contents

1. The resting screen — header, KPI band, chips
2. The sidebar — the tool's conscience
3. Anatomy of one answer card
- 4–12. Nine sample questions with real outputs
13. The refusals — two questions it will not answer
14. What changes when the AI is on
15. The demo cheat sheet


1 • The resting screen — header, KPI band, chips

Launch with `./run.sh` (or `docker compose up`). Before any question, the page already presents the quarter. Walk it top to bottom:

The title and promise

"**Meridian Pipeline Intelligence**" with one sentence under it: *every answer shows its source records, its formula, and its assumptions — numbers are computed, never generated.* The product's whole contract, in the first two lines. The meta line beneath anchors time: **Q1 2026 final • Q2 2026 in progress, day 32 of 91 (as of May 02, 2026) • 8 audited analyses.**

The four KPI cards (they count up from \$0 when the page loads)



POINT AT IT AND SAY

Before anyone asks a question, the tool has already disclosed two things a dashboard wouldn't: most of that \$518K was arguably counted last quarter — see "strictly new \$145,000" — and even winning every open deal leaves a \$484,000 hole. The quarter's real headline is: create pipeline now.

The five suggested chips

One row of capsules — **Enterprise vs quota • Reps at risk • Q2 vs Q1 pace • Regional view • Data quality.** Hovering shows the full question each one sends. Below them, the chat box: *"Ask about quota, reps, segments, regions, or data quality..."*

2 · The sidebar — the tool's conscience

Everything in the left rail is the tool being honest about itself. Top to bottom:

Data status

"Data validated · 5 reconciliations passed" with a green dot, then the facts: as of May 02 2026 · Q1 (final) 88 deals · Q2 snapshot 92 deals · Reps 10 · Day 32 of 91, and a progress bar at 35%. If any reconciliation had failed, this page would show an error instead of numbers — the app fails closed.

AI assist

A toggle — **"LLM routing & narration"** — and the contract line: *numbers always computed, never generated*, plus which provider is active. Flip it off live any time: the app keeps answering in its own wording. This switch is your proof that AI is a convenience here, not a dependency.

Observability

Live counts from the audit log: questions logged, answered vs refused, AI summaries verified, response times — and a "Recent questions" list. Every answer the tool has ever given is on record in `logs/queries.jsonl`.

Known data issues

The five findings as compact rows — F1 reused deal IDs (\$1.97M of pipeline), F2 closed deals reopened, F3 overdue deals, F4 revenue counted in both quarters (\$373K), F5 edited history. Detected automatically at load; fix the files and these rows disappear.

Methodology

An expandable panel with every metric definition, the six assumptions, and the reconciliation list — the whole rulebook, one click from every answer.

POINT AT IT AND SAY

The sidebar is the part I'm proudest of, and it never changes: the tool's own health checks, its own known problems, and its own rulebook — permanently on screen. Nothing about this product asks to be trusted; it shows its working.

3 · Anatomy of one answer card

Every answer has the same seven parts, in the same order. Learn this once and every sample below reads instantly:

1. **The routing line** (top-left, small caps) — what the question was interpreted as, and what it matched on: a suggested question, a keyword, a rep's name, or *"LLM classification, validated against the approved intent catalog."* Misunderstandings are visible here, never hidden.
2. **The trust badge** (top-right) —  **Verified**  **Verified with assumptions** or  **Insufficient data** . Fixed rules, not feelings.
3. **The answer** — two to four sentences. With AI on, it's the model's phrasing that survived validation; otherwise the engine's own wording. Below it, the verification line: *"AI summary · N numbers verified against computed results."*
4. **Metric tiles** — the key figures as cards.
5. **A chart**, when one helps.
6. **Data-quality warnings** — always visible, orange-edged, quantified in dollars.
7. **Three tabs** — **Source data** (the actual rows, with problem rows tagged), **Calculation** (numerator, denominator, filters, thresholds, written out), **Assumptions**.

Each new answer also fires a small toast, bottom-right: *"Answer verified — N numbers checked against the engine."*

4 · Sample question 1 — the Enterprise chip

Click: the *Enterprise vs quota* chip (sends "How is the Enterprise segment tracking against quota this quarter?").

ENTERPRISE Q2 QUOTA TRACKING · SUGGESTED QUESTION

● Insufficient data

Enterprise has closed \$165,000 of its \$3,650,000 Q2 quota (4.5%) with 35% of the quarter elapsed. Open pipeline of \$3,600,000 covers the remaining gap 1.03x — below the 1.28x needed at Q1's realized win rate. Caution: \$165,000 of that booked figure was already recorded as won in the final Q1 book — strictly new Q2 revenue is \$0.

CLOSED-WON
\$165,000
1 deal, Apr 1 – May 2

Q2 QUOTA
\$3,650,000
4 Enterprise reps

ATTAINMENT
4.5%
vs 35% elapsed

OPEN PIPELINE
\$3,600,000
20 open deals

COVERAGE
1.03x
break-even 1.28x

\$165,000 of Enterprise's Q2 booked revenue (100%) comes from deals recorded as Closed Won in the final Q1 book whose close dates moved into Q2 (F4). Counting only strictly new wins, Enterprise Q2 bookings are \$0.

\$1,475,000 of Enterprise open pipeline sits on reused deal IDs (F1); treat pipeline-based conclusions with care.

\$305,000 of Enterprise open pipeline is deals already closed in the final Q1 book that have reopened (F2).

Source data: the numerator (1 row — Sunridge Systems) · the quota basis (4 reps) · the open pipeline (20 rows, problem rows tagged)

POINT AT IT AND SAY

Notice the badge is red — the tool is downgrading its own answer. The 4.5% is arithmetically perfect and strategically misleading, and it tells you why: its entire numerator is one deal your closed Q1 books already counted. Open Source data — the numerator is one visible row. When the answer is one deal, you get shown the deal.

5 · Sample question 2 — the Reps at risk chip

Click: *Reps at risk* (sends "Which reps are at risk of missing Q2?").

REPS AT RISK OF MISSING Q2 · SUGGESTED QUESTION

● Verified with assumptions

3 of 10 reps are At risk — behind pace with pipeline mathematically short of their remaining quota: Kevin Marsh (REP-06), Lisa Park (REP-09), Ryan Cole (REP-10). 4 are on Watch, and 3 cannot be responsibly classified because most of their pipeline sits on reused deal IDs.

AT RISK

3

pipeline < remaining quota

WATCH

4

behind on one signal

ON TRACK

0

INSUFFICIENT DATA

3

>50% reused-ID pipeline

\$373,000 of booked Q2 revenue was already recorded as won in the final Q1 book (F4), concentrated in: REP-01 (\$165,000 — 100% of their booked); REP-08 (\$208,000 — 71%).

Day 32 of 91: 7 of 10 reps have \$0 booked. Q1 showed the same early pattern (21% at day 32) and finished at 102% — read "At risk" as needs pipeline scrutiny, not a verdict.

Chart: coverage per rep with two reference lines — 1.0x (gap unclosable below this) and 1.28x (break-even at Q1's real win rate).
Source data: the full 10-rep table with every input to the rules.

POINT AT IT AND SAY

Three verdicts here, and the third is the one that matters: the tool refuses to grade Patel, Torres, and Williams — over half their pipeline sits on corrupted records. It reports them as a data problem, not a performance problem. That refusal protects careers, and it's a judgment a hallucinating tool cannot make.

6 · Sample question 3 — the Q2 vs Q1 chip

Click: Q2 vs Q1 pace (sends "How does Q2 attainment compare to where we were at the same point in Q1?").

Q2 VS Q1 AT THE SAME POINT · SUGGESTED QUESTION

● Verified with assumptions

Q2 is running behind Q1's pace: \$518,000 closed (8.4% of quota) through day 32, versus \$1,260,000 (21.3%) at the same point in Q1 — about 41% of Q1's dollar pace. Q1 went on to finish at 102.0%. Counting only strictly new Q2 wins (\$145,000), the gap is wider still.

Q2 BOOKED (DAY 32) \$518,000 5 deals · 8.4%	Q1 BOOKED (DAY 32) \$1,260,000 10 deals · 21.3%	GAP −12.9% pts equal days elapsed	STRICTLY-NEW Q2 \$145,000 excluding F4	Q1 FINAL 102.0% \$6,041,000 · back-loaded
--	--	--	---	--

Q2's window has one more business day (23 vs 22) and one more calendar day (91 vs 90); both favor Q2 — so the gap is not an alignment artifact.

The Q1 side is a reconstruction: the final Q1 file filtered to close dates on or before Feb 01.

Chart: two attainment bars, Q1 gray vs Q2 blue. Source data: the 5 Q2 wins and the 10 Q1 wins, row by row.

POINT AT IT AND SAY

The fairness question is answered before anyone asks it: we aligned on equal days elapsed, checked the calendar both ways, and both residual effects favor Q2 — so the 41%-of-pace finding is real. Balanced by honesty: Q1 also started slow and finished at 102%. The conclusion isn't panic; it's create pipeline now.

7 • Sample question 4 — the Regional view chip

Click: *Regional view* (sends "How is each region performing this quarter?").

REGIONAL PERFORMANCE • KEYWORD "REGION"

● Verified with assumptions

By deal region, Southeast leads Q2 bookings with \$208,000; total Q2 bookings are \$518,000 against a \$6,200,000 quota, with \$5,198,000 of open pipeline distributed across four regions.

CENTRAL

\$145,000

quota \$1,330,000 • pipe \$1,830,000

NORTHEAST

\$0

quota \$1,780,000 • pipe \$1,110,000

SOUTHEAST

\$208,000

quota \$1,370,000 • pipe \$1,315,000

WEST

\$165,000

quota \$1,720,000 • pipe \$943,000

12 active deals (\$1,335,000) sit in a region different from their owning rep's home region. Revenue follows the deal's region while quota follows the rep's home region — regional attainment is approximate; treat it as directional.

Source data: the per-region table, plus the 12 cross-region deals listed row by row — the source of the approximation.

POINT AT IT AND SAY

A small example of the product's character: it could have printed regional attainment percentages that look precise. Instead it tells you the quota basis and the deal basis don't perfectly align — 12 deals, \$1.3 million — and labels the view directional. An approximate number labeled approximate beats a precise-looking number that quietly isn't.

8 · Sample question 5 — the Data quality chip

Click: *Data quality* (sends "What's wrong with my data?").

DATA QUALITY REPORT · KEYWORDS "DATA" + QUALITY

● Verified

Five distinct data-integrity findings affect 28 of 92 deals in the current snapshot. The most consequential: \$373,000 of Q2 "bookings" was already recorded as Q1 revenue (F4), and \$1,970,000 of open pipeline sits on reused deal IDs (F1). Every affected answer in this tool is flagged automatically; fixing the source data clears the flags.

F1

13 deals

\$1,970,000 at stake

F2

7 deals

\$576,000 at stake

F3

2 deals

\$87,000 at stake

F4

3 deals

\$373,000 at stake

F5

3 deals

\$260,000 at stake

Chart: red bars of value touched per finding. Source data: all 28 affected rows tagged per finding, plus the final-Q1-book versions of the re-dated and edited deals for side-by-side comparison.

POINT AT IT AND SAY

Note the badge — this is the one answer that's fully green. The single thing this dataset supports with total confidence is a precise description of its own problems. "What's wrong with my data" is a one-click question on purpose: data quality stopped being an IT topic the moment it started deciding which reps look like failures.

9 · Sample question 6 — type "How is Priya Patel doing?"

INDIVIDUAL REP PERFORMANCE · REP NAME "PRIYA PATEL"

● Insufficient data

Priya Patel (REP-03, Enterprise Northeast) has closed \$0 of a \$1,000,000 Q2 quota (0.0%), with \$740,000 open pipeline covering the gap 0.74x. Status: Insufficient data. This rep's pipeline is dominated by reused deal IDs — the records need correction before performance conclusions are reliable.

QUOTA

\$1,000,000

BOOKED

\$0

0 deals

ATTAINMENT

0.0%

pace 0.00

PIPELINE

\$740,000

4 open deals

COVERAGE

0.74x

STATUS

Insufficient data

61% of this rep's open pipeline value sits on reused deal IDs (F1).

POINT AT IT AND SAY

On paper she looks terrible — zero booked, pipeline short. A naive tool would flag her as failing. This one refuses: 61% of her pipeline is corrupted records, so the verdict is "fix the data before judging the person." The action item is a records cleanup, then re-ask — and the tool will grade her the moment the data deserves it.

10 · Sample question 7 — type "How is Tom Bradley doing?"

The deliberate contrast with Priya — the best performer, with an asterisk the tool refuses to hide.

INDIVIDUAL REP PERFORMANCE · REP NAME "TOM BRADLEY"

● Insufficient data

Tom Bradley (REP-08, Mid-Market Southeast) has closed \$293,000 of a \$520,000 Q2 quota (56.3%), with \$220,000 open pipeline covering the gap 0.97x. Status: Watch. Caution: \$208,000 of this booked figure was already recorded as won in the final Q1 book (re-dated close date) — strictly new Q2 revenue is \$85,000.

QUOTA \$520,000	BOOKED \$293,000 3 deals	ATTAINMENT 56.3% pace 1.60	PIPELINE \$220,000 3 open deals	COVERAGE 0.97x	STATUS Watch
---------------------------------------	---	---	--	--------------------------------------	------------------------------------

\$208,000 of this rep's Q2 booked revenue (71%) is from F4 re-dated deals also counted in the final Q1 book.

39% of this rep's open pipeline value sits on reused deal IDs (F1).

POINT AT IT AND SAY

Here's the tool being fair in the other direction. Tom is the quarter's standout — 56.3%, pace 1.60 — and his risk status is still Watch, computed on the numbers as recorded. But the trust badge goes red anyway, because 71% of his booked figure is re-dated Q1 revenue: strictly new, he's at \$85,000. Same rules for the struggling rep and the star — that symmetry is what fairness looks like in software.

11 · Sample question 8 — type "How is SMB tracking?"

SEGMENT PERFORMANCE · SEGMENT "SMB"

● Verified with assumptions

SMB has closed \$0 of its \$510,000 Q2 quota (0.0%) with 35% of the quarter elapsed. Open pipeline of \$413,000 covers the remaining gap 0.81x — below the 1.28x needed at Q1's realized win rate.

BOOKED

\$0

0 deals

QUOTA

\$510,000

2 SMB reps

ATTAINMENT

0.0%

PIPELINE

\$413,000

11 open deals

COVERAGE

0.81x

\$96,000 of SMB open pipeline is deals already closed in the final Q1 book that have reopened (F2).

2 open deals in this segment are past their expected close date (OPP-066, OPP-075).

POINT AT IT AND SAY

Worth a sentence in passing: SMB has booked nothing — but it's also the only segment whose pipeline is completely free of reused-ID records. Zero revenue, cleanest data. The badges track data health, not performance — which is exactly what a trust grade should do.

12 · Sample question 9 — type "Show me the pipeline"

COMPANY PIPELINE SUMMARY · KEYWORD "PIPELINE"

● Verified with assumptions

Company-wide, Q2 stands at \$518,000 closed of a \$6,200,000 quota (8.4%) with \$5,198,000 of open in-quarter pipeline (46 deals). Even if every open deal closed, \$484,000 of quota would remain uncovered — new pipeline creation is the immediate priority.

Q2 QUOTA

\$6,200,000

all 10 reps

BOOKED

\$518,000

5 deals · 8.4%

OPEN PIPELINE

\$5,198,000

46 deals

UNCOVERED

\$484,000

even at 100% win rate

\$373,000 of the \$518,000 booked (72%) comes from F4 re-dated deals — strictly new Q2 revenue is \$145,000.

\$1,970,000 (38%) of open pipeline sits on reused deal IDs (F1).

\$466,000 of open pipeline is deals already closed in the final Q1 book and reopened in Q2 (F2).

Chart: pipeline by stage — Discovery \$1,678,000 · Proposal \$1,977,000 · Negotiation \$1,543,000 · **Prospecting \$0 — empty**. Source data: the 10 largest open deals (reused-ID rows tagged).

POINT AT IT AND SAY

Look at the stage chart: the first stage — Prospecting — is completely empty. Nothing new is being seeded. This quarter's problem is coverage; next quarter's problem is being created right now. That empty bar might be the most strategic pixel in the product.

13 · The refusals — two questions it will not answer

Type: "What will Q3 revenue be?"

NOT SUPPORTED · OUT OF SCOPE: NO DATA FOR "Q3"

● Refused

I cannot answer that reliably from the current dataset. The available data supports questions about quota, attainment, deals, representatives, segments, regions, and quarter-over-quarter comparisons.

Below the refusal: the list of the eight supported analyses. The refusal is also written to the audit log.

Type: "Who will win the World Cup?" (or invite the panel to ask anything)

NOT SUPPORTED · NO APPROVED INTENT MATCHED

● Refused

Same honest refusal — no improvisation, no numbers, no guess.

POINT AT IT AND SAY

This is the moment to hand the keyboard to the panel: "ask it anything — genuinely." Off-menu questions hit an honest refusal; future questions hit "no data exists." The previous tool's fatal habit was improvising when it didn't know. This one treats "I can't answer that reliably" as a first-class answer — and every refusal becomes backlog for the next audited analysis.

14 · What changes when the AI is on

Everything above is the deterministic engine — it's what you get with the sidebar toggle off, and it's the exact material the AI is allowed to rephrase. With the toggle on, three things change (and only three):

1. Conversational phrasings route correctly

Type **"Are we going to make the number this quarter?"** — keyword rules alone can't place that, so the model picks from the menu. The routing line shows it: *"Company pipeline summary · LLM classification, validated against the approved intent catalog."* Same numbers as chapter 12 — the model only chose the menu item.

2. Answers read more naturally — with a proof line

The headline becomes the model's phrasing of the same computed package, and beneath it: *"AI summary · 11 numbers verified against computed results · deterministic wording under Calculation."* If the model had included even one number the engine didn't produce — even "\$3.65M" as shorthand — the entire summary would be rejected, the engine's wording shown instead, and the rejection displayed. The guardrail is visible either way.

3. The toast confirms each check

Bottom-right after every answer: *"Answer verified — N numbers checked against the engine."*

POINT AT IT AND SAY

And if the AI provider goes down mid-demo? Flip the toggle — or don't even bother: the app degrades to exactly what you've seen in this document, automatically. We watched that happen during a real provider outage. The AI is the polish; the trust is in the machine underneath.

15 · The demo cheat sheet

#	DO	THE ONE THING TO LAND
0	Resting screen (1 min)	"\$484,000 uncovered even at a 100% win rate — the quarter's real headline, before any question."
1	Enterprise chip	Red badge; numerator is one row; strictly-new = \$0.
2	Reps at risk chip	Three at risk — and three the tool refuses to judge. Data problem ≠ performance problem.
3	Q2 vs Q1 chip	41% of last quarter's pace, and the alignment favors Q2 — the gap is real.
4	Type "Are we going to make the number this quarter?"	LLM routing line + "N numbers verified" — AI on a visible leash.
5	Type "How is Priya Patel doing?"	The refusal to grade a person on corrupted records.
6	Data quality chip	The only green answer; "fix the data and the flags clear themselves."
7	Type "What will Q3 revenue be?" — or hand the panel the keyboard	Honest refusal. "Eight trusted answers beat eight thousand plausible ones."

Recovery: AI down → sidebar toggle off, narrate graceful degradation as a feature. App down → this document and the one-page leave-behind carry the entire walkthrough on paper.