


Meridian Pipeline Intelligence

A trust-first answer engine for pipeline questions — every answer carries its source records, its formula, and its assumptions. Numbers are computed, never generated.
Q1 2026 final · Q2 2026 in progress, day 32 of 91 (as of May 2, 2026) · 8 audited analyses · Streamlit + pandas, LLM-optional (Claude / Gemini interchangeable)

TRUST ARCHITECTURE — THE MODEL CHOOSES FROM A MENU; IT NEVER COOKS



Worst case is a **visible** misroute the user can correct — never an invisible wrong number. Six intents shipped day one; capability grows intent-by-intent, each audited before it answers a leader.

THE THREE REQUIRED ANSWERS — AS COMPUTED

Q · Enterprise vs Q2 quota ● **Insufficient data**

4.5% · \$165,000 of \$3,650,000

Coverage 1.03x vs 1.28x break-even (derived from Q1's realized 77.9% win rate). Flagged: the entire numerator is re-dated Q1 revenue — strictly-new Enterprise bookings are **\$0**.

Q · Reps at risk of missing Q2 ● **Verified w/ assumptions**

3 at risk · 4 watch · 3 unclassifiable

Marsh, Park, Cole: behind pace *and* pipeline < remaining quota. Patel, Torres, Williams: 54–67% of pipeline on reused deal IDs — reported as a data problem, not a performance problem.

Q · Q2 vs Q1, same point (day 32) ● **Verified w/ assumptions**

8.4% vs 21.3% — \$518K vs \$1,260K

Q2 runs at ~41% of Q1's dollar pace; strictly-new basis is \$145,000 — the gap only widens. Context: Q1 was also back-loaded and finished at 102% (\$6,041,000).

DATA-INTEGRITY FINDINGS — DETECTED AT LOAD, NEVER HARDCODED

FINDING	SCALE	POLICY
F1 Deal IDs reused for different companies across snapshots (OPP-076... 088)	13 deals · \$1.97M 38% of open pipeline	Included as current system of record; excluded from all deal-history comparisons; flagged on every touched answer
F2 Deals closed in the final Q1 book reopened or re-lost in Q2	7 deals · \$576K	Q1 reporting follows the closed book; Q1-revenue/Q2-pipeline overlap disclosed
F3 Open deals past their expected close date	2 deals · \$87K	Kept in pipeline; hygiene flag
F4 Closed-Won in both snapshots, close date moved Q1 → Q2 — same revenue counted in both quarters (OPP-042, 060, 074)	3 deals · \$373K 72% of "Q2 booked"	Both figures shown everywhere; answers dominated by it downgrade their own trust badge to red
F5 Closed records edited after the Q1 book closed (e.g. value cut \$155K → \$120K)	3 deals · \$260K	Q1 follows the Q1 file; edits surfaced, not silently adopted

KEY DEFINITIONS & ASSUMPTIONS

- Attainment** = Closed-Won ÷ quota. Open pipeline reported separately, never blended.

Coverage = open pipeline ÷ remaining quota; 1.28x break-even derived from Q1's realized value win-rate, not an industry heuristic.

Same point = equal days elapsed (day 32: May 2 ↔ Feb 1); residual calendar effects favor Q2, so the gap is real.

Risk = fixed published thresholds on pace, coverage, and record reliability — not a predictive model. Weighted forecasting deliberately excluded: no stage-outcome history exists to calibrate probabilities honestly.

As-of date (May 2, 2026) from the brief — the files carry no snapshot timestamp.

VERIFICATION

- 38 automated tests, golden numbers hand-checked**

5 reconciliations at every load, fail-closed

Routing eval **22/24** on golden set (rules + LLM)

Narration validator: **0** uncomputed numbers can render

Adversarial audit found **F4 after** tests passed — detection now automated

WHAT IT PROVES / DOESN'T

- Proves:** trustworthy plain-English answers over governed metrics are buildable in take-home scope, with hallucination structurally impossible in the numeric path.

Doesn't: production scale, live-CRM messiness, open-ended NL. Next: read-only Salesforce/Snowflake with snapshot journaling — making F4-style re-dating impossible to hide, not just detectable.