

The Mock Interview

Twenty-two questions calibrated to your interviewer — a former natural-language-BI founder who now leads Deloitte's Palantir Foundry & AIP practice — each with a speakable first-person answer and a coaching note. Practice until the answers are yours, not the page's.

Candidate: Sai Saran · Artifact: Meridian Pipeline Intelligence · Interviewer profile: AI & Data Solutions AVP, NL-BI founder (datavoice), human-centered design facilitator, MSDS w/ data-ethics focus

How to use this (three passes)

1. **Pass 1 — read silently**, whole document. Note which answers feel foreign.
2. **Pass 2 — answer aloud** reading only the QUESTION, then compare with the script. Don't memorize wording — memorize the *skeleton* (each answer is: claim → evidence from the artifact → honest limit).
3. **Pass 3 — blind rapid-fire**: cover everything, answer from the question list in Chapter 10, 30 seconds each.

Rule for every answer: one claim, one concrete number or demo moment as evidence, one honest limitation. Never bluff Palantir specifics — the fallback line is in Chapter 3, Q3.

1 • Round 0 — The opening (they always start here)

INTERVIEWER

"So — walk us through what you built."

YOUR ANSWER — first person, ~45–75 seconds

Meridian's problem wasn't missing data — it was that leaders stopped asking questions: answers were slow, and the one AI tool they tried invented numbers in front of the CCO. So I built a question-answering system with receipts. A leader asks in plain English; the question is routed to one of eight audited analyses; a deterministic engine computes the answer — no AI can touch a number — and the answer ships with its source rows, its formula, and a trust badge that downgrades itself when the data can't support it. Along the way the tool's own checks found five real problems in the sample data, including \$373,000 of revenue counted in both quarters — which is why its very first answer grades itself red instead of pretending. Let me show you the client's own three test questions live. *[open the app]*

COACH • Keep this to 90 seconds and get to the demo — this interviewer builds products; he'd rather click than listen. The run-order: Enterprise (red badge) → Reps at risk → Q2 vs Q1 → type "Are we going to make the number this quarter?" (LLM tier) → "How is Ryan Ferrera doing?" (refusal — the smile beat) → "How is Ryan Cole doing?" (the real Ryan, answered) → Data quality (green).

INTERVIEWER

"What did you cut, and what would you do with another four hours?"

YOUR ANSWER — first person, ~45–75 seconds

Every cut passed one test: does it trade verifiability for capability? Free-form ask-anything — cut, because an invented analysis can be invisibly wrong, which is the failure you already lived. Weighted forecasting — cut, because the data has no stage history, so any win probabilities would be invented numbers. A separate API and auth — real, but theater on sample data; they're week one of production. With four more hours: a "challenge this number" button that files a review ticket with full context, product-line analyses, and one-click PDF export of any answer with its evidence and flags.

COACH • The four-hours answer should be **boringly concrete** — that's what makes it credible. Resist naming ten things.

2 · Round 1 — NL-to-BI depth (the datavoice founder's home turf — expect the sharpest follow-ups here)

INTERVIEWER

"Two VPs ask 'what's our attainment?' and expect two different numbers. What does your tool do?"

YOUR ANSWER — first person, ~45–75 seconds

It gives both of them the same number with the definition printed on the answer: Closed-Won revenue over quota, open pipeline never blended, dates and filters written out in the Calculation tab. That doesn't settle whose definition is right — that's a governance decision — but it ends the meeting where two leaders quote different numbers with the same word. And when sales ops disagrees with the definition, it's a one-line config change, versioned, and every historical answer in the audit log still shows which definition it used. In this dataset that scenario is real: "Q2 booked" is \$518K by the system of record and \$145K strictly-new — the tool shows both and tells you there's a decision to make.

COACH · This is the **semantic-layer question** in disguise. The phrase "definition printed on the answer" is your anchor.

INTERVIEWER

"How do you handle a genuinely ambiguous question?"

YOUR ANSWER — first person, ~45–75 seconds

Three behaviors, all visible. First, transparent keyword rules take the unambiguous cases. Second, the LLM tier classifies the conversational ones — but it can only choose off the menu, its choice is re-validated, and the interpretation is printed at the top of the answer: "interpreted as company pipeline summary." Third, if nothing fits, it refuses and lists what it can answer. And I don't assume routing works — I measure it: 22 of 24 on a golden phrasing set, and the important property is that every miss fails to a visible refusal or a labeled interpretation, never to a wrong number.

COACH · If he pushes "but the interpretation could still be wrong" — agree instantly: **"yes — visibly wrong, which the user corrects in one click. I traded invisible errors for visible ones."**

INTERVIEWER

"Tools like this demo beautifully and are abandoned by month two. Why is yours different?"

YOUR ANSWER — first person, ~45–75 seconds

Because the thing that kills them is the first wrong answer — one hallucination and the leader never comes back. This tool is built so its first wrong answer is a refusal or a red badge, not a lie. Beyond that, adoption is designed in: every refusal is logged and becomes the backlog for the next audited analysis, so the tool grows toward what leaders actually ask; and the pilot's success metric is behavioral — are leaders asking questions again — because that was the original disease. I'd also say honestly: eight intents is a narrow start. Narrow and trusted is survivable; broad and wrong is not.

COACH · He lived this with datavoice. If it feels natural, this is the moment for: **"you built datavoice — I'd genuinely love to hear what killed adoption there, and whether this approach fixes it or repeats it."** Once, humbly.

INTERVIEWER

"Why a chat interface at all? Why not just better dashboards?"

YOUR ANSWER — first person, ~45–75 seconds

Because the client's own words were "leaders stop asking." A dashboard answers the questions someone predicted last quarter; the CCO's next question is never on it, so it routes to an analyst and dies in a queue. Chat removes the translation layer — but naked chat is the thing that burned them. So the design is: the front door of a bot, the engine room of a dashboard, plus self-awareness neither has — badges, refusals, receipts. And the channel is swappable: everything renders from one answer contract, so the same engine could sit in Slack, a portal, or voice.

COACH · "Front door of a bot, engine room of a dashboard" is your most quotable line. Land it once, don't reuse it.

3 · Round 2 — Trust, governance, and the Palantir lens

INTERVIEWER

"Where exactly can the model lie to me in your system?"

YOUR ANSWER — first person, ~45–75 seconds

In numbers — nowhere. There's no code path from model output into arithmetic: the engine imports no LLM library, and the model's two jobs are choosing from the intent menu and phrasing computed results, where a validator extracts every number from its prose and vetoes the whole sentence if even one isn't in the computed packet — including sneaky ones, like abbreviating \$3,650,000 to "\$3.65M." Where it can still mislead is wording: a correct-numbers sentence with wrong emphasis. I mitigate that by always keeping the deterministic wording one tab away and showing the verification count on screen — but I won't claim that risk is zero.

COACH · The unprompted honest limit at the end is what separates an 8 from a 10 on this question.

INTERVIEWER

"You've essentially built a tiny semantic layer with governed functions. Did you realize you rebuilt the Foundry/AIP pattern?"

YOUR ANSWER — first person, ~45–75 seconds

Deliberately so — at take-home scale. Eight audited functions over a governed metric layer, an LLM allowed only to select among them and phrase their outputs, validation on both sides, and an audit log underneath. That's the AIP philosophy: the model gets leverage through governed actions, never through raw access. I haven't shipped on Foundry itself — I want to be straight about that — but I built this architecture because I believe it's the only NL-analytics pattern that survives contact with a skeptical enterprise, and it's why the Palantir side of your practice genuinely interests me.

COACH · The most important answer in this document. Say "I haven't shipped on Foundry" before he asks — he leads the practice and will smell any stretch. The last clause turns it into motivation.

INTERVIEWER

"Your trust badges — who decided the rules, and can I audit them?"

YOUR ANSWER — first person, ~45–75 seconds

Fixed rules, all in one config file an auditor can read in a minute: for example, if more than half of a booked figure is re-dated revenue, the answer downgrades itself to red; if more than half a rep's pipeline sits on reused IDs, the tool refuses to classify the person at all. The thresholds that could be arbitrary aren't — the 1.28× coverage break-even is derived from Q1's actual 77.9% win rate, their own data. And the same rules apply symmetrically: my best performer, Tom Bradley at 56% attainment, still carries a red badge because 71% of his booked is re-dated. The star gets the same scrutiny as the struggler.

COACH · The Bradley symmetry is your proof the rules aren't vibes. Have his numbers cold: \$293K booked, \$208K re-dated, \$85K strictly new.

4 · Round 3 — Human-centered design

INTERVIEWER

"Who is the user — and what did you deliberately NOT build for them?"

YOUR ANSWER — first person, ~45–75 seconds

The user is a sales leader with thirty seconds, not an analyst with an afternoon — that drove everything. For them I did NOT build: query builders, filter panels, or free-form asking; the KPI band answers the pre-question glance, five chips cover the questions they actually ask, and there's an assistant-only mode that hides everything except the conversation. The analyst isn't ignored — they're served by the receipts: every answer carries its rows and formula, so the analyst's verification job takes thirty seconds instead of an afternoon too. Two users, one artifact, different layers.

COACH · He trains people on user-centered design — "two users, one artifact, different layers" is his language.

INTERVIEWER

"If Meridian were a real client, how would you have run discovery?"

YOUR ANSWER — first person, ~45–75 seconds

A visioning session with the CCO's leadership team — but anchored in artifacts, not opinions: bring the last ten questions each leader actually sent to an analyst, and the last report they didn't trust and why. Those ten questions become the initial intent menu — not my guess at eight. Then a data-forensics pass before any build, because in this exercise the forensics changed everything: five findings, including revenue counted in two quarters. And once live, discovery never stops — every refusal the tool logs is a leader telling you what to build next.

COACH · "Discovery never stops — the refusal log is the backlog" is the senior move here.

INTERVIEWER

"Tell me about a design decision you reversed."

YOUR ANSWER — first person, ~45–75 seconds

The landing page. My first version restated the whole engagement brief — cards, stats, the story. It was thorough and it was wrong: overloaded, text-heavy, built for me rather than the user. I rebuilt it as an assistant-home — a greeting, one input, four example questions — and moved the brief into a collapsed "about" section. The general lesson I took: my instinct under pressure is to prove work by showing volume; the discipline is to prove it by showing restraint. The same reversal happened with the sidebar — five paragraphs of data findings became five one-line rows with details on demand.

COACH · Reversals answered honestly build more trust than pretending the first draft was right. Keep it light.

5 · Round 4 — Scale & production

INTERVIEWER

"Two million rows in Snowflake, nulls everywhere, fields sales ops repurposed years ago. What survives from your prototype?"

YOUR ANSWER — first person, ~45–75 seconds

The contract survives; the implementation swaps. Every answer is already a structured packet — answer, evidence, calculation, flags — so the eight analyses become governed, versioned warehouse queries; pandas doesn't come along. The validation posture survives and moves upstream: my five load-time reconciliations become pipeline gates that block bad data before it's queryable. And the single most important production upgrade is snapshot journaling — recording every change to every deal — because it converts my biggest finding from detectable-by-audit to impossible-to-hide: closed-won dates moving across a quarter boundary becomes a standing alert, not a forensic discovery. What honestly doesn't survive: I haven't proven performance at that scale, and real CRM mess will surface defect classes I haven't conceived — that's why the posture is detect, disclose, downgrade, rather than a fixed list of checks.

COACH · He architects this for a living. The phrase "detectable → impossible to hide" is the one to land.

INTERVIEWER

"We deploy into government environments. What breaks in yours?"

YOUR ANSWER — first person, ~45–75 seconds

Less than you'd expect, by accident of the trust design. The LLM is optional — the app degrades to fully deterministic operation with no external calls, which is air-gap-friendly. The container builds with the test suite as a gate, everything is auditable down to a query-by-query log, and there's no data leaving the boundary except the optional narration call — which is the first thing I'd swap for an in-boundary model. What's genuinely missing for gov: access control, classification handling, and accreditation paperwork — real work, but additive to this architecture rather than fighting it.

COACH · Don't oversell clearance-world knowledge; the honest "missing" list is the credibility.

6 · Round 5 — Business & strategy

INTERVIEWER

"Forget the tool. You're in front of Meridian's CCO Monday morning — what do you tell her to do?"

YOUR ANSWER — first person, ~45–75 seconds

Three things, in order. First, create pipeline this week: even winning every open deal leaves \$484,000 of quota uncovered, and the Prospecting stage is empty — nothing new is being seeded, so next quarter's problem is being created right now. Second, run a CRM hygiene sprint before the next QBR: twenty-eight records carry defects, three of her reps can't even be evaluated because of them, and \$373,000 of revenue currently lives in two quarters at once — that last one needs finance in the room, not just sales ops. Third, until journaling makes re-dating impossible, run the quarter on two numbers side by side: booked \$518K, strictly-new \$145K. They're different enough to change decisions.

COACH · Numbers from memory, no hedging — this answer IS the FDE job. If you nail one answer tomorrow, make it this one.

INTERVIEWER

"How does this prototype become an engagement?"

YOUR ANSWER — first person, ~45–75 seconds

The wedge is the data findings — they're already worth the engagement fee, and they create urgency the demo alone wouldn't. Phase one: three-leader pilot on read-only replicated CRM data, six weeks, two commitments — every refused question becomes intent backlog, every challenged number gets an answer within a day. Success metric is behavioral: are leaders asking questions again. Phase two grows the intent library from observed demand and moves the checks into the pipeline. I'd keep one governance rule non-negotiable: no free-form query path without a separate risk sign-off — the constrained surface is why it's trustworthy, and it's always the first thing someone asks to remove.

COACH · "The findings are the wedge" is strategy-speak he'll recognize. The non-negotiable at the end shows spine.

7 · Round 6 — Ethics & behavioral (his MSDS focus area was data ethics — expect at least one)

INTERVIEWER

"Your tool labels human beings 'at risk.' Talk to me about that."

YOUR ANSWER — first person, ~45–75 seconds

It's the most ethically loaded feature in the product, and it shaped three decisions. The rules are published, not predictive — pace, coverage, record reliability — so a rep can see exactly why they're flagged and what changes it. The tool refuses to grade anyone whose records are majority-corrupted — three of the ten reps today — because "bad data" must never silently become "bad performance"; that's a career-affecting distinction. And symmetrically, the top performer's inflated numbers get flagged too. For production I'd add governance before rollout: HR and sales ops decide who sees rep-level labels, because a risk flag that leaks into comp conversations without context is harm, not insight.

COACH · This is the answer where you can visibly care. Slow down for "bad data must never silently become bad performance."

INTERVIEWER

"Tell me about a time you were wrong during this project."

YOUR ANSWER — first person, ~45–75 seconds

The best one: my first working version passed every test I'd written — and was still materially wrong. It reported Q2 revenue at \$518K without knowing that 72% of it was Q1 revenue with quietly moved close dates. My tests were green because I'd only tested for defects I'd imagined. I found it by adversarially auditing my own outputs against the raw files — trying to refute my own numbers — and then automated the detection so it's now a standing check with its own tests. It permanently changed my claim about the tool: not "it can't be wrong," but "being wrong is detectable, attributable, and cheap to fix."

COACH · Tell it as a story with a turn, not a confession. This is your strongest three minutes — don't rush it.

INTERVIEWER

"How much of this did AI build, honestly?"

YOUR ANSWER — first person, ~45–75 seconds

A lot of the code — I used an AI coding assistant as a pair-programmer throughout, which your brief explicitly encourages. What I owned: the read of the client problem as trust rather than features, the architecture decision to keep the model out of the numeric path, every metric definition and threshold, the adversarial audit that caught the double-counted revenue after tests passed, and the verification standard — 42 tests, five reconciliations, a measured router. I can defend any line of it, and I'd point out the recursion: I used AI with guardrails to build a tool whose entire thesis is AI with guardrails. How I used it is, I think, the demonstration the brief was asking for.

COACH · Zero defensiveness. Confidence here comes from the fact that you *can* defend any line — which is why Pass 3 practice matters.

8 · Round 7 — The required question from the brief

INTERVIEWER (VERBATIM FROM THE BRIEF)

"The last AI tool this client used lost executive trust because it hallucinated. How does your approach avoid that problem — and where does it still fall short?"

YOUR ANSWER — first person, ~45–75 seconds

It avoids it structurally, not behaviorally. The old tool's model both understood questions and produced numbers, so every answer was a potential fabrication. I separated the jobs permanently: the model only chooses from eight audited analyses and phrases their computed results — and a validator vetoes any sentence containing a number the engine didn't produce. There is no code path for a hallucinated figure to reach the screen; it can ignore every instruction and still can't lie in numbers. Where it still falls short, honestly: wording can mislead even when numbers are right; routing can pick a wrong menu item — visibly, but wrongly; the thresholds are calibrated on one quarter; and the deepest limit — deterministic checks only catch defects a human conceived. I know that last one is real because it happened to me: my audit found \$373K of double-counted revenue after every test was green. That's why my claim isn't "never wrong" — it's "wrong is visible, traceable, and cheap to fix," which I'd argue is the only trust claim an honest engineer can make.

COACH · This is the one answer worth knowing nearly verbatim — the brief literally announces it's coming. 90 seconds, unhurried.

9 · Curveballs — short answers, 20 seconds each

CURVEBALL	YOUR 20-SECOND ANSWER
"Ask it something ridiculous right now."	Hand over the keyboard. Refusal path is safe. "That refusal is the feature — the last tool improvised."
"Why Streamlit and not a real front end?"	Prototype speed with a disposable UI; all logic lives in the engine, which is already API-shaped. The UI is the cheapest thing to replace.
"Why no RAG? Everyone does RAG."	RAG grounds answers in documents; this data is tables, where retrieval should be an exact filter, not a similarity guess. The day they bring contracts, RAG becomes audited intent number nine.
"Is 8.4% attainment bad?"	Against 35% elapsed it looks bad; against Q1's 21%-at-day-32 → 102%-finish pattern it's early. The bad number is coverage below 1.28× — that one doesn't fix itself.
"Which number is real — \$518K or \$145K?"	Both computed correctly from different premises; choosing is a finance-policy decision. My job was making sure they know there's a decision.
"What if Gemini goes down mid-demo?"	Flip the toggle — everything still works. Happened to me live during a provider billing outage; users saw plainer wording, zero errors.
"How long did this really take?"	The deterministic core fits the 1–2h scope; I chose to extend it — LLM layer, audit, deployment, docs — to show the production direction, and I can point at which parts belong to which phase.
"What's your biggest technical debt?"	Keyword routing brittleness and single-quarter threshold calibration. Both have written fixes in the roadmap.

10 · Blind rapid-fire list (Pass 3 — cover everything above)

1. Walk me through what you built.
2. What did you cut and why?
3. Two VPs expect different "attainment" numbers — what happens?
4. How do you handle ambiguity?
5. Why won't this be abandoned like every other chat-BI tool?
6. Why chat instead of dashboards?
7. Where can the model lie to you?
8. You rebuilt the AIP pattern — did you know?
9. Who audits your trust badges?
10. Who's the user and what didn't you build?
11. How would you run discovery?
12. A design decision you reversed?
13. What survives at 2M rows?
14. What breaks in a government environment?
15. What should the CCO do Monday?
16. How does this become an engagement?
17. Your tool labels people "at risk" — defend that.
18. A time you were wrong in this project?
19. How much did AI build?
20. The hallucination question (brief, verbatim).
21. Which is real: \$518K or \$145K?
22. Why no RAG?

Your questions for him (pick two)

- "In your AIP work, where do you draw the line between governed functions and letting the model compose freely — and has that line moved this year?"
- "When you built datavoice, what actually killed adoption — and does my constrained-menu approach fix it or repeat it?"
- "What does the best FDE you've worked with do in week one at a client?"

11 · Self-scoring & final checks

After each practice pass, score yourself 1–5

DIMENSION	YOU'RE A 5 WHEN...
Numbers fluency	\$518K/\$145K/\$373K/\$484K/1.28x/77.9%/22-of-24 come out without pausing
Honest limits	Every answer volunteers a limitation before being asked
Story beats	F4 story and Bradley symmetry told as 60-second narratives with a turn
Palantir humility	"I haven't shipped on Foundry" said plainly, followed by the pattern argument
Brevity	No answer over 90 seconds unless invited deeper

Night-before checklist

- ☐ Run the app once end-to-end (`./run.sh`), ask all five chips + the Ryan pair
- ☐ Close every unrelated browser tab; share window, not desktop
- ☐ PDFs open in tabs: architecture · leave-behind · guardrails
- ☐ Rehearse Round 8 (hallucination answer) aloud twice
- ☐ Rehearse the CCO-Monday answer aloud twice
- ☐ Sleep. A rested 80% beats an exhausted 100%.